

Fast Regression of the Tritium Breeding Ratio in Tokamak Fusion Reactors

G Van Goffrier¹ and P Mánek^{1,2}, V Gopakumar³, N Nikolau¹,
J Shimwell³, I Waldmann¹

¹ Department of Physics and Astronomy, University College London, Gower Street, London WC1E 6BT, UK

² Institute of Experimental and Applied Physics, Czech Technical University, Husova 240/5, Prague 110 00, Czech Republic

³ UK Atomic Energy Authority, Culham Science Centre, OX14 3DB Abingdon, UK

E-mail: `graham.vangoffrier.19@ucl.ac.uk`, `petr.manek.19@ucl.ac.uk`

Abstract. The tritium breeding ratio (TBR) is an essential quantity for the design of modern and next-generation Tokamak nuclear fusion reactors. Representing the ratio between tritium fuel generated in breeding blankets and fuel consumed during reactor runtime, the TBR depends on reactor geometry and material properties in a complex manner. In this work, we explored the training of surrogate models to produce a cheap but high-quality approximation for a Monte Carlo TBR model in use at the UK Atomic Energy Authority. We investigated possibilities for dimensional reduction of its feature space, reviewed 9 families of surrogate models for potential applicability, and performed hyperparameter optimisation. Here we present the performance and scaling properties of these models, the fastest of which, an artificial neural network, demonstrated $R^2 = 0.985$ and a mean prediction time of $0.898\,\mu\text{s}$, representing a relative speedup of $8 \cdot 10^6$ with respect to the expensive MC model. We further present a novel adaptive sampling algorithm, Quality-Adaptive Surrogate Sampling, capable of interfacing with any of the individually studied surrogates. Our preliminary testing on a toy TBR theory has demonstrated the efficacy of this algorithm for accelerating the surrogate modelling process.

Index terms— magnetic moment, solar neutrinos, astrophysics Submitted to: *J.*

Phys. G: Nucl. Part. Phys.

1. Introduction

The analysis of massive datasets has become a necessary component of virtually all technical fields, as well as the social and humanistic sciences, in recent years. Given that rapid improvements in sensing and processing hardware have gone hand in hand with the data explosion, it is unsurprising that software for the generation and interpretation of this data has also attained a new frontier in complexity. In particular, simulation procedures such as Monte Carlo (MC) event generation can perform physics predictions even for theoretical regimes which are not analytically tractable. The bottleneck for such procedures, as is often the case, lies in the computational time and power which they necessitate.

Surrogate models, or metamodels, can resolve this limitation by replacing a resource-expensive procedure with a much cheaper approximation [?]. They are especially useful in applications where numerous evaluations of an expensive procedure are required over the same or similar domains, e.g. in the parameter optimisation of a theoretical model. The term “metamodel” proves especially meaningful in this case, when the surrogate model approximates a computational process which is itself a model for a (perhaps unknown) physical process [?]. There exists a spectrum between “physical” surrogates which are constructed with some contextual knowledge in hand, and “empirical” surrogates which are derived purely from samples of the underlying expensive model.

In this project, in coordination with the UK Atomic Energy Authority (UKAEA), we sought to develop a surrogate model for the tritium breeding ratio (TBR) in a Tokamak nuclear fusion reactor. Our expensive model was an MC-based neutronics simulation, *Paramak*[‡], which returns a prediction of the TBR for a given configuration of a spherical Tokamak. We took an empirical approach to the construction of this surrogate, and no results described here are explicitly dependent on prior physics knowledge.

For the remainder of Section 1, we will define the TBR and set the context of this work within the goals of the UKAEA. In Section 2 we will describe our datasets generated from the expensive model for training and validation purposes, and the dimensionality reduction methods employed to develop our understanding of the parameter domain. In Section 3 we will present our methodologies for the comparison testing of a wide variety of surrogate modelling techniques, as well as a novel adaptive sampling procedure suited to this application. After delivering the results of these approaches in Section 4, we will give our final conclusions and recommendations for further work.

[‡] Provided by collaborator Jonathan Shimwell, at UKAEA.

This section was copied from the report “as is” and needs to be shortened.

Figure 1. Typical single-null reactor configuration as specified by BLUEPRINT [?]:
1 — plasma, 2 — breeding blankets

1.1. Problem Description

Nuclear fusion technology relies on the production and containment of an extremely hot and dense plasma. In this environment, by design similar to that of a star, hydrogen atoms attain energies sufficient to overcome their usual electrostatic repulsion and fuse to form helium [?]. Early prototype reactors made use of the deuterium (^2H , or D) isotope of hydrogen in order to achieve fusion under more accessible conditions, but lead to limited success. The current frontier generation of fusion reactors, such as the Joint European Torus (JET) and the under-construction International Thermonuclear Experimental Reactor (ITER), make use of tritium (^3H , or T) fuel for further efficiency gain. Experimentation at JET dating back to 1997 [?] has made significant headway in validating deuterium-tritium (D-T) operations and constraining the technology which will be employed in ITER in a scaled-up form.

However, tritium is much less readily available as a fuel source than deuterium. While at least one deuterium atom occurs for every 5000 molecules of naturally-sourced water, and may be easily distilled, tritium is extremely rare in nature. It may be produced indirectly through irradiation of heavy water (D_2O) during nuclear fission, but only at very low rates which could never sustain industrial-scale fusion power.

Instead, modern D-T reactors rely on tritium breeding blankets, specialised layers of material which partially line the reactor and produce tritium upon neutron bombardment, e.g. by



where T represents tritium and ${}^7\text{Li}$, ${}^6\text{Li}$ are the more and less frequently occurring isotopes of lithium, respectively. ${}^6\text{Li}$ has the greatest tritium breeding cross-section of all tested isotopes [?], but due to magnetohydrodynamic instability of liquid lithium in the reactor environment, a variety of solid lithium compounds are preferred.

The TBR is defined as the ratio between tritium generation in the breeding blanket per unit time and tritium fuel consumption in the reactor. The MC neutronics simulations previously mentioned therefore must account for both the internal plasma dynamics of the fusion reactor and the resultant interactions of neutrons with breeding blanket materials. Neutron paths are traced through a CAD model (e.g. Figure 1) of a reactor with modifiable geometry.

The input parameters of the computationally-expensive TBR model therefore fall into two classes. Continuous parameters, including material thicknesses and packing ratios, describe the geometry of a given reactor configuration. Discrete categorical parameters further specify all relevant material sections, including coolants, armours, and neutron multipliers. One notable exception is the enrichment ratio, a continuous

parameter denoting the presence of ^6Li . Our challenge, put simply, was to produce a fast TBR function which takes these same input parameters and approximates the MC TBR model with the greatest achievable regression performance.

2. Methodology

Assuming that data is appropriately treated to eliminate redundant features, we proceed to propose surrogate models and criteria used for their evaluation. The task all presented surrogates strive to solve can be formulated using the language of conventional regression problems. In this section, we focus on interpretation in the scheme of supervised learning with decoupled and adaptive sampling.

Labeling the expensive MC TBR model $f(x)$, a surrogate is a mapping $\hat{f}(x)$ that yields similar images as $f(x)$. In other words, $f(x)$ and $\hat{f}(x)$ minimise a selected dissimilarity metric. Furthermore, in order to be considered *viable*, surrogates are required to achieve expected evaluation time lower than that of $f(x)$.

In the decoupled sampling approach that is further described in Section 2.2, we first gather a sufficiently large training set of samples $\mathcal{T} = \{(x^{(i)}, f(x^{(i)}))\}_{i=1}^N$ to describe the behaviour of $f(x)$ across its domain. Depending on a specific model family and appropriate choice of its hyperparameters, surrogate models $\hat{f}(x)$ are trained to minimise the empirical risk $R_{\text{emp.}}$ with respect to \mathcal{T} and a model-specific loss function \mathcal{L} , where the empirical risk is defined as $R_{\text{emp.}}(\hat{f} | \mathcal{T}, \mathcal{L}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(\hat{f}(x^{(i)}), f(x^{(i)}))$.

The adaptive sampling approach that we characterise in Section 2.3 can be viewed as a more general problem. Rather than fixing the training set \mathcal{T} for the entire duration of training, multiple sets $\{\mathcal{T}_k\}_{k=0}^K$ are used, such that $\mathcal{T}_{k-1} \subset \mathcal{T}_k$ for all $k > 1$. The first set \mathcal{T}_0 is initialised randomly to provide a *burn-in*, and is repeatedly extended in epochs, whereby each epoch trains a new surrogate $\hat{f}_k(x)$ on \mathcal{T}_k using the supervised learning procedure, evaluates its performance, and forms a new set \mathcal{T}_{k+1} by adding more samples to \mathcal{T}_k . This permits the learning algorithm to condition the selection of new samples in \mathcal{T}_{k+1} on the evaluation results of $\hat{f}_k(x)$ in order to maximise improvement of surrogate performance over complex regions within the feature space.

2.1. Metrics

Aiming to provide objective comparison of a diverse set of surrogate model families, we define a multitude of metrics to be tracked during experiments. Following the motivation of this work, two desirable properties of surrogates arise: (a) their capability to approximate the TBR MC model well and (b) their prediction time. An ideal surrogate would maximise the former while minimising the latter.

Table 1 provides an exhaustive listing and description of metrics recorded in our experiments. For regression performance analysis, we include a selection of absolute metrics to assess the approximation capability of surrogates, and set practical bounds on the expected uncertainty of their predictions. In addition, we also track relative

measures that are better-suited for comparison between this work and others, as they maintain invariance with respect to the selected domain and image space. For complexity analysis, surrogates are assessed in terms of wall time§ elapsed during training and prediction. This is motivated by common practical use cases of our work, where models are trained and used as drop-in replacements for the expensive MC TBR model. Since training set sizes remain to be determined, all times are reported per a single datapoint. Even though some surrogates support acceleration by means of parallelisation, sequential processing of samples was ensured to achieve comparability between considered models.||

Regression performance metrics	Mathematical formulation / description	Ideal value [units]	
Mean absolute error (MAE)	$\sum_{i=1}^N y^{(i)} - \hat{y}^{(i)} /N$	0	[TBR]
Standard error of regression S	$\text{StdDev}_{i=1}^N \{ y^{(i)} - \hat{y}^{(i)} \}$	0	[TBR]
Coefficient of determination R^2	$1 - \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2 / \sum_{i=1}^N (y^{(i)} - \bar{y})^2$	1	[rel.]
Adjusted R^2	$1 - (1 - R^2)(N - 1)/(N - P - 1)$	1	[rel.]
Complexity metrics			
Mean training time $\bar{t}_{\text{trn.}}$	(wall training time of $\hat{f}(x)$)/ N_0	0	[ms]
Mean prediction time $\bar{t}_{\text{pred.}}$	(wall prediction time of $\hat{f}(x)$)/ N	0	[ms]
Relative speedup ω	(wall evaluation time of $f(x)$)/($N\bar{t}_{\text{pred.}}$)	$\rightarrow \infty$	[rel.]

Table 1. Metrics recorded in supervised learning experiments. In formulations, we work with a training set of size N_0 and a test set of size N , TBR values $y^{(i)} = f(x^{(i)})$ and $\hat{y}^{(i)} = \hat{f}(x^{(i)})$ denote images of the i th testing sample in the expensive model and the surrogate respectively. Furthermore, the mean $\bar{y} = \sum_{i=1}^N y^{(i)}/N$ and P is the number of input features.

To prevent undesirable bias in results due to training set selection, all metrics are collected in the scheme of k -fold cross-validation with a conventional choice of $k = 5$. In this setting, a sample set is uniformly divided into 5 disjoint folds, each of which is used as a test set for models trained on the remaining 4.¶ Having repeated the same experiment for each such run, the overall value of individual metrics is reported in terms of their mean and standard deviation over all folds.

2.2. Decoupled Sampling

In our experiments, we evaluate and compare surrogates in effort to optimise against metrics described in Section 2.1. To attain meaningful and practically usable results, we require a sufficiently large and diverse pool of surrogate families to review. This is described by the listing in Table 2. The presented selection of models

§ Real time elapsed during computation measured by a chronometer, here by means of the Python `time` package.

|| The only exception to this are artificial neural networks, which require considerable amount of processing power for training on conventional CPU architectures.

¶ Unless explicitly stated otherwise, we use 1 fold for testing and the 4 remaining folds for training. This gives 80% to 20% train-test ratio.

includes basic techniques suitable for linear regression enhanced by the kernel trick or dimension lifting, methods driven by decision trees, instance-based learning models, ensemble regressors, randomised algorithms, artificial neural networks and mathematical approaches developed specifically for the purposes of surrogate modelling. For each of these families, a state-of-the-art implementation was selected and adapted to operate with TBR samples.

Surrogate	Acronym	Implementation	Hyperparameters
Support vector machines [?]	SVM	SciKit Learn [?]	3
Gradient boosted trees [?, ?, ?]	GBT	SciKit Learn	11
Extremely randomised trees [?]	ERT	SciKit Learn	7
AdaBoosted decision trees [?]	ABT	SciKit Learn	3
Gaussian process regression [?]	GPR	SciKit Learn	2
k nearest neighbours	KNN	SciKit Learn	3
Artificial neural networks	ANN	Keras (TensorFlow) [?]	2
Inverse distance weighing [?]	IDW	SMT [?]	1
Radial basis functions	RBF	SMT	3

Table 2. Considered surrogate model families.

While some of the presented model families are clearly determined by an explicit choice of a learning algorithm, others may vary considerably depending on hyperparameter configuration. A good example of the latter group are artificial neural networks, which in addition to conventional hyperparameters enable to control network architecture through selection from various parametric graph templates (illustrated in Figure 2). This allows us to realise a simplistic network architecture search during the hyperparameter tuning procedure.

Placeholder

Figure 2. Selected parametric neural network architectures. All layers except the last use ReLU activation. Prediction information flow is indicated by arrows.

2.2.1. Experiments The presented surrogate candidates are evaluated in four experimental cases:

- (i) Hyperparameter tuning in a simplified domain.
- (ii) Hyperparameter tuning in full domain.
- (iii) Scaling benchmark.
- (iv) Model comparison.

The aim of the initial experiments is to use a relatively small subset of collected TBR samples to determine hyperparameters of considered surrogates. Since this process requires learning the behaviour of an unknown, possibly expensive mapping – here a function that assigns cross-validated metrics to a point in the hyperparameter

domain – it in many aspects mirrors the primary task of this work with the notable extension of added utility to optimise. In order to avoid undesirable exponential slowdown in exhaustive searches of a possibly high-dimensional parameter space, Bayesian optimisation [?] is employed as a standard hyperparameter tuning algorithm. We set its objective to maximise R^2 and perform 1000 iterations.⁺

In the first experiment, efforts are made to maximise the possibility of success in surrogates that are prone to suboptimal performance in discontinuous spaces. This follows the notion that, if desired, performance of such models may be replicated by training separate instances to model each continuous subregion of the domain independently. To this end, data are limited to a single slice from run 2, and discrete features are completely withheld from evaluated surrogates. This is repeated for each of the four available slices to investigate variance in behaviour under different discrete feature assignments. The second experiment conventionally measures surrogate performance on the full feature space. Here, in extension of the previous case, surrogates work with samples comprised of discrete as well as continuous features.

The objective of the last two experiments is to exploit the information gathered by hyperparameter tuning. In the third experiment, the 20 best-performing hyperparameter configurations of each family (with respect to R^2) are used to perform training on progressively larger sets to investigate their scaling properties. Following that, the fourth experiment aims to produce surrogates suitable for practical use by retraining selected well-scaling instances on large training sets to satisfy the goals of this work.

2.3. Adaptive Sampling

Placeholder

Figure 3. Schematic of QASS algorithm

All of the surrogate modelling techniques studied in this project face a common challenge: their accuracy is limited by the quantity of training samples which are available from the expensive MC TBR model. Adaptive sampling procedures can improve upon this limitation by taking advantage of statistical information which is accumulated during the training of any surrogate model. Rather than training the surrogate on a single sample set generated according to a fixed strategy, sample locations are chosen periodically during training so as to best suit the model under consideration.

Adaptive sampling techniques appear frequently in the literature and have been specialised for surrogate modelling. Garud’s [?] “Smart Sampling Algorithm” achieved notable success by incorporating surrogate quality and crowding distance scoring to identify optimal new samples, but was only tested on a single-parameter domain. We

⁺ Hyperparameter tuning of each surrogate family was terminated after 2 days. Instances that reached this limit may be identified in Table ?? in the Appendix.

theorised that a nondeterministic sample generation approach, built around Markov Chain Monte Carlo methods (MCMC), would fare better for high-dimensional models by more thoroughly exploring all local optima in the feature space. MCMC produces a progressive chain of sample points, each drawn according to the same symmetric proposal distribution* from the prior point. These sample points will converge to a desired posterior distribution, so long as the acceptance probability for these draws has a particular functional dependence on that posterior value (see [?] for a review).

Many researchers have embedded surrogate methods into MCMC strategies for parameter optimisation [?, ?], in particular the ASMO-PODE algorithm [?] which makes use of MCMC-based adaptive sampling to attain greater surrogate precision around prospective optima. Our novel approach draws inspiration from ASMO-PODE, but instead uses MCMC to generate samples which increase surrogate precision throughout the entire parameter space.

We designed the Quality-Adaptive Surrogate Sampling algorithm (QASS, Figure 3) to iteratively increment the training/test set with sample points which maximise surrogate error and minimise a crowding distance metric (CDM) [?] in feature space. On each iteration following an initial training of the surrogate on N uniformly random samples, the surrogate was trained and absolute error calculated. MCMC was then performed on the error function generated by performing nearest-neighbor interpolation on these test error points. The resultant samples were culled by 50% according to the CDM, and then the n highest-error candidates were selected for reintegration with the training/test set, beginning another training epoch. Validation was also performed during each iteration on independent, uniformly-random sample sets.

3. Results

4. Conclusion

5. Acknowledgements

6. References

* An adaptive MCMC procedure [?], which adjusts an ellipsoidal proposal distribution to fit the posterior, was also implemented but not fully tested.